Data Platform Design

Document date: 2024-07-15 Classification: Internal

Table of Contents

1. Ingestion 2
2. Orchestration 3
3. Governance 4

1. Ingestion

Events flow from Kafka into Snowflake via the analytics pipeline. Staging mirrors production topology but runs at roughly one tenth of the capacity. Remember to tag pull requests with the owning team so reviews get routed correctly. We should revisit the alert thresholds after the next game day. Load tests should be coordinated in the platform channel so they do not trip alerts.

2. Orchestration

Airflow schedules batch jobs. On-call handover happens every Monday at 10:00 local time. Postmortems are blameless and are published within five business days. Flaky integration tests in CI are tracked on the quality board; please do not just re-run them. Remember to tag pull requests with the owning team so reviews get routed correctly.

3. Governance

Database schema migrations must be backward compatible for at least one release. Structured JSON logging is mandatory; free-text logs are dropped by the log pipeline. Tracing uses OpenTelemetry with a ten percent head-based sampling rate. Service-level objectives are reviewed quarterly with product owners. Capacity planning for the holiday peak starts in September every year.