Platform Architecture Overview

Document date: 2024-02-10 Classification: Internal

Table of Contents

1. Topology 2
2. Request Flow 3
3. Caching 4
4. Observability 5

1. Topology

The platform runs 12 core services on Kubernetes (EKS) in eu-west-1. On-call handover happens every Monday at 10:00 local time. The dashboards were updated to show saturation alongside latency and errors. On-call handover happens every Monday at 10:00 local time. Remember to tag pull requests with the owning team so reviews get routed correctly. Secrets must never be committed; the pre-commit hook scans for high-entropy strings.

2. Request Flow

External traffic enters through the API Gateway, which validates JWTs and applies rate limits before routing to backend services. Postmortems are blameless and are published within five business days. Postmortems are blameless and are published within five business days. Tracing uses OpenTelemetry with a ten percent head-based sampling rate. Staging mirrors production topology but runs at roughly one tenth of the capacity.

3. Caching

Redis is used by checkout-service for carts and by api-gateway for rate-limit counters. Service-level objectives are reviewed quarterly with product owners. Structured JSON logging is mandatory; free-text logs are dropped by the log pipeline. Staging mirrors production topology but runs at roughly one tenth of the capacity.

4. Observability

Feature work is paused during incident follow-up weeks until action items are closed. Service-level objectives are reviewed quarterly with product owners. Every service has a README that links to its dashboards and runbook. Capacity planning for the holiday peak starts in September every year. Infrastructure changes go through Terraform plans that need two approvals. Load tests should be coordinated in the platform channel so they do not trip alerts.